

ECSE 426 - Microprocessor Systems
Lab Report 1: Analog Data Acquisition, Filtering, and
Digital I/O

Harley Wiltzer (260690006)
Matthew Lesko

February 19, 2018

Contents

0.1	Implementation	2
0.1.1	The User Input Module	2
0.1.2	The Data Acquisition Module	2
0.1.3	The Data Processing Module	4
0.2	The Output Module	6
0.3	Testing and Observations	8
A	Circular Lists	10

List of Figures

1	Debouncing button presses	3
2	Flow of ADC processing	5
3	Selecting 7 segment display with transistors	7

List of Tables

1	Accuracy of ADC by resolution	3
2	Extracting digits from floating point number	7
3	Results from testing the voltmeter with various input voltages	8

0.1 Implementation

The design of the voltmeter was fairly complex and was composed of several modules. These modules included the [user input module](#) for processing user input via the push button, the [data acquisition module](#) for digitizing analog data on the board, the [data processing module](#) for filtering the data and associating meaning to its digital values, and the [output module](#) for displaying the data to the user. This section will be divided into several subsections, each corresponding to a module, in order to organize the design decisions that were made.

Since all of these modules needed to work together, they had to be synchronized appropriately, and this was achieved with the `SysTick` timer. The `SysTick` timer invokes interrupts at a chosen frequency, and the interrupt handler was used to coordinate all of the modules. Since the configuration parameters of the `SysTick` timer were heavily influenced by the modules described above, they will be explained independently in the sections following where the design decisions were made.

0.1.1 The User Input Module

One of the requirements of the voltmeter was to provide three display modes to the user: a display of the RMS voltage, and a display for each of minimum and maximum voltage updated within the past ten seconds. As such, there was a need for user input to switch between these display modes. This was achieved with the user button on the STM32F407 board, which allowed the user to cycle through each of the display modes by pressing the button.

The first challenge dealt with how to process the button presses. There were two main options: polling for button presses and handling interrupts. Since polling the button at every iteration in a loop seemed inefficient, the interrupt method was chosen. Therefore, an NVIC interrupt was configured at priority 0 for `EXTI0`. However, even with the interrupts set up, there was still a major challenge to correctly process the button presses, as one button press often caused several interrupts. This could have been due to button bouncing, or possibly the fact that a what seems like a short *click* to a human actually goes over several clock cycles of the processor. To prevent this from occurring, it was decided to enforce a time delay between consecutive button press handling routines, and this was achieved using the `SysTick` timer. The process is shown in [Figure 1](#).

In brevarium, the `EXTI0_IRQHandler()` function (invoked by the button press interrupt) asserts a `change_mode` signal, and when the `SysTick_Handler()` function (invoked by `SysTick` interrupt) sees that, it waits for 32 consecutive `SysTick` interrupts before taking action. The number 32 was achieved via trial and error, as it was unknown exactly how long an average human button press lasts. That being said, this number was chosen at a `SysTick` frequency of 200Hz, so a delay of 160ms was imposed.

0.1.2 The Data Acquisition Module

The data acquisition module was responsible for gathering analog data and digitizing it so it could be processed. Firstly, however, it was helpful to set up a digital to analog converter (DAC) in order to test the performance of the analog to digital converter (ADC). Setting up the DAC was fairly straightforward, and most of the work was carried out by the HAL Cube software. The DAC

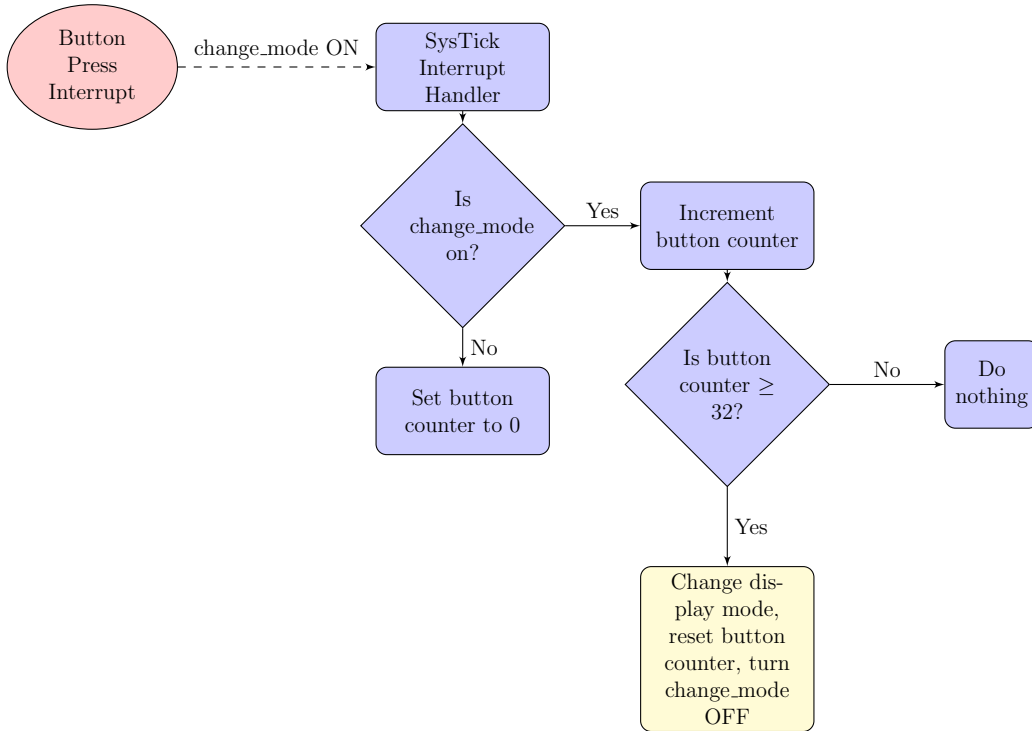


Figure 1: Debouncing button presses

was configured on channel 1, and it wrote to pin PA4 on the board. Furthermore, its resolution had to be chosen. Since the performance of the voltmeter ultimately depended on the resolution of the ADC, the resolution of the DAC was chosen to be the same as that of the ADC, which was 8 bits, right aligned. This decision will be explained when discussing the ADC parameters below. Finally, the DAC needed to output some analog voltage. A value was chosen arbitrarily and passed to the DAC via the `HAL_DAC_SetValue()` driver function, and the conversion was instantiated via `HAL_DAC_Start()`. This starts a conversion in polling mode, which was deemed appropriate for the purposes of this experiment as the conversion would only occur once.

Setting up the ADC was considerably more complicated. Again, the basic initialization was done by the HAL Cube software, and the ADC1 unit was set up on channel 1. Single conversion mode was chosen, as it was required for one conversion to occur at a given frequency. Once again, the resolution had to be determined. Since it was known that the displayed voltages would be shown with two decimal places of precision, the user could only see voltages in increments of 0.01V. The

Table 1: Accuracy of ADC by resolution

Resolution (bits)	Voltage difference between consecutive digital values (V)
6	0.047
8	0.012
10	0.003
12	4.89e-4

accuracy of the ADC by its resolution is shown in [Table 1](#). The accuracy is defined here as the change in voltage when increasing the digital reading by 1. Since the voltage range of the ADC is 3V, the accuracies were calculated according to

$$A = \frac{3}{2^R} \quad (1)$$

where A is the accuracy (rightmost column) and R is the resolution (leftmost column). Clearly, a resolution of 6 bits is not a great choice, as it cannot resolve voltages within 0.047V from each other. Since the display of the voltmeter allowed two decimal places, this accuracy is insufficient. With a resolution of 10 bits, however, the accuracy is relatively high. At an accuracy of 0.003V, the display would only change after a change of 4 in the digital reading of the ADC. The 8 bit resolution could resolve voltages that are 0.012V apart which is very close to the accuracy on the display. Ultimately, the 8 bit and 10 bit resolutions were the main contenders, because the 8 bit resolution is slightly worse than that of the display, and the 10 bit resolution is much stronger than that of the display. In the end, the 8 bit resolution was chosen as it was deemed strong enough for the purposes of this experiment, it would cause lower power consumption, and it matched one of the possible resolutions of the DAC which made the code simpler.

Next, the conversion mode of the ADC had to be chosen. Polling mode was not considered a viable option, since ADC conversions would happen frequently and thus polling would waste a considerable portion of the CPU's cycles. Although DMA was a very good alternative, the developers did not have time to do the requisite research. Therefore, interrupt mode was selected. Despite the conversion mode, however, the frequency of ADC conversions remained to be implemented. A sample rate of 50Hz was required, so the `SysTick` interrupts were used to time the ADC conversions. Since the `SysTick` interrupts were occurring at 200Hz (see the [Output subsection](#) below), it was required to implement a prescaler in the `SysTick_Handler()` function in order to sample at 50Hz. The process of handling ADC conversions is shown in [Figure 2](#). The `SysTick_Handler()` maintains a `counter` and increments it at every `SysTick` interrupt, that is to say, at 200Hz. Since ADC conversions were to be taken at 50Hz, they had to be triggered at a rate four times less frequent than `SysTick`. Thus, the `HAL_ADC_StartIT()` function was called on every fourth `SysTick` interrupt to start ADC conversions at 50Hz. This was accomplished by starting the ADC conversion when the `counter` variable was a multiple of 4.

0.1.3 The Data Processing Module

Given the digital readings from the ADC, the next step was to filter them to reduce noise and to then translate the readings into meaningful values. These steps were carried out in the `ConvCpltCallback()` function (see [Figure 2](#)). In order to reduce noise, the digital readings were passed through a 4th order FIR filter, which effectively output the average over the past 5 inputs. A 9th order filter was experimented (which effectively output the average over the past 10 inputs), but it had a negligible effect on the performance of the voltmeter. Since the 9th order filter required considerably more memory, the 4th order filter was chosen. By passing each new ADC reading through the filter, potential noise will be “smoothed out”, as it will be converted into the average of that reading and the previous 4.

Upon filtering the ADC readings, the filtered output was then translated into its analog rep-

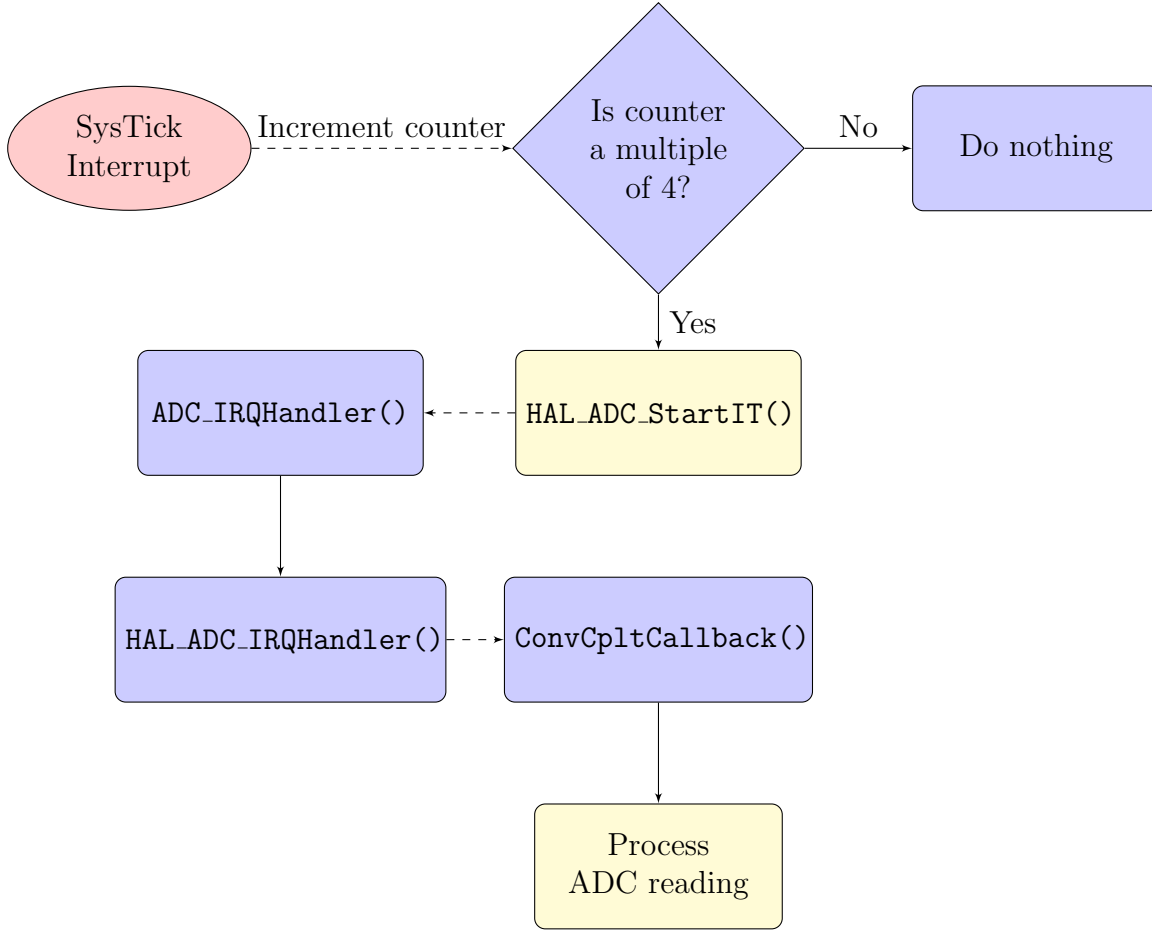


Figure 2: Flow of ADC processing

resentation according to

$$x_a = V_{DD} \left(\frac{\mathcal{F}(x_d)}{2^R - 1} \right) \quad (2)$$

where x_a is the analog representation of the new ADC reading, $\mathcal{F}(x_d)$ is the filtered digital ADC reading, R is the resolution (in bits) of the ADC, and V_{DD} is the highest voltage that the ADC is rated for. Finally, the analog value is passed to the `plot_point()` function, which is responsible for curating the inputs into the form in which they should be displayed.

The goal of the `plot_point()` function is to maintain the values of the RMS voltage over the past 10 seconds, as well as the minimum and maximum voltages over the past 10 seconds, while making efficient use of memory. It was decided to design this in a “moving window” fashion, meaning the values to be displayed will always taken into account data from within the past 10 seconds, and not in discrete 10 second blocks. With the requirement of keeping memory usage down to at most 20 samples, this was a difficult task. Firstly, it was decided that the minimum and maximum voltage do not need to be updated very frequently. Of the 20 samples, 5 samples were allocated to each of minimum and maximum voltages, and the remaining 10 samples were reserved for RMS voltage. With only 5 samples for each of minimum and maximum voltages, each sample would have to represent the minimum or maximum sample in a 2 second interval. Likewise, for RMS, each stored sample represents the running RMS over a 1 second interval. The minimum samples,

maximum samples, and RMS samples were each stored in their own circular list (see [Appendix A](#)) for efficient insertion and removal. For each 100 samples (2 seconds) passed to `plot_point()`, the minimum and maximum were calculated and added to their respective circular lists. For the RMS calculations, a new value was added to the RMS circular list for every 50 samples (1 second), however the value to be added could not be the RMS exactly. RMS is calculated as follows

$$\text{RMS} = \sqrt{\frac{1}{|N|} \sum_{n \in N} n^2} \quad (3)$$

where N is the set of samples to calculate the RMS over. Since the square root operation is non-linear, the RMS itself cannot be passed to the RMS circular list every second. Rather, the *sum of squares* $\sum_{n \in N} n^2$ is passed, and the RMS is calculated over the past 10 sum of squares values.

With the circular lists mentioned above, the `plot_point()` function was able to return the RMS, minimum, and maximum voltages over the past 10 seconds fairly easily. For the minimum and maximum voltages, this consisted of simply finding the minimum value in the minimum voltages circular list and the maximum value in the maximum values circular list. The RMS over the past 10 seconds was computed as follows:

$$\text{RMS}_{10} = \sqrt{\frac{1}{50|S|} \sum_{s \in S} s} \quad (4)$$

where S is the set of the 10 most recent second-long sum of squares calculations described above. Since each second consists of 50 readings, the sum of sums of squares had to be divided by $50|S|$, because it summed over 50 points for each $s \in S$.

Note that although the RMS is only being updated once per second and the minimum and maximum are only being updated once every 2 seconds, the circular lists still store enough data to report the exact RMS, minimum, and maximum over a moving 10 second window. Due to the constraint on memory usage, the update frequency of the RMS, minimum, and maximum could not be improved with a moving window design.

0.2 The Output Module

The output module was responsible for displaying the data that the [data processing module](#) computed using LEDs and 7 segment displays on the STM32F407 development board. This required setting up the GPIO configuration for three LEDs to indicate the display mode as well as 11 output pins to control the 7 segment displays. All LEDs and output pins were configured in push-pull mode at low speed with `GPIO_NOPULL` set for the `Pull` parameter. These configurations were made because it was desired to output LOW and HIGH voltages instead of HIGH and high impedance. Resistors were included on the breadboard for the purpose of current limiting.

Controlling the output LEDs for the display mode was very simple. Since the `SysTick` interrupt handler was responsible for updating a `display_mode` variable upon button presses, the output LED was set by resetting all three LEDs and then setting the one corresponding to the appropriate display mode in the program's main loop.

Manipulating the set of 7 segment displays was more difficult. In order to save wiring, the display took 8 data inputs (for each of the segments on a given display) and four selection inputs to select with display to illuminate. The strategy therefore was to multiplex the displays and update the data lines depending on which display was currently being illuminated. Assuming the the clock governing the switching of the display being illuminated is fast enough, the human eye should not be able to see the displays turning on and off; rather, it should look like all displays are on simultaneously. The selection inputs of the display chip were each connected to a common cathode line for a different 7 segment display. To turn on one of the 7 segment displays, the common cathode should be grounded. To turn it off, the common cathode should be disconnected altogether. To accomplish this, 3 output pins on the board were designated to the selection lines (the fourth 7 segment display was not being used). These output pins were then connected to their respective common cathodes through NPN transistors as follows:




Figure 3: Selecting 7 segment display with transistors

ing the current flowing through it. When the digit select pin is HIGH, the transistor is active and a logic LOW is connected to the common cathode. However, when the digit select pin is LOW, the transistor is off, and from the common cathode's perspective, it sees an open circuit. Therefore, these transistor circuits can be used to turn on and off the different 7 segment displays.

In order to display a number, first the appropriate number is selected according to the display mode. In the main program loop, the digit select pins are set according to the counter variable maintained by `SysTick`, such that each display is on once every three `SysTick` interrupts. Depending on which display is on at a given moment, the appropriate digit is extracted to be sent to the display according to Table 2. Given the digit calculated according to Table 2, it is then necessary

Table 2: Extracting digits from floating point number

Digit to be displayed	Function to extract appropriate digit
Ones	$x\%10$
Tenths	$(10x)\%10$
Hundredths	$(100x)\%10$

to compute compute the 8 data bits sent to the 7 segment display. This was done by manually associating segments to digits and making a function returning a byte containing which segments are on or off. With this byte, the output pins on the board are set HIGH or LOW accordingly, thereby displaying the numbers on the 7 segment display.

However, one design challenged still remained. It was imperative to ensure that the frequency

of the displays being multiplexed was fast enough such that it would look like all displays look simultaneously on to the human eye. Originally, the `SysTick` frequency was set to 50Hz to accommodate the ADC, however at 50Hz each display turned on and off at $\frac{50}{3}$ Hz, and the displays were noticeably flickering. Therefore, the `SysTick` timer had to be modified to reduce the flickering. Frequencies that were multiples of 50Hz were attempted in order to ease the prescaling for the ADC samples, and the lowest frequency that removed any sight of flickering was found to be 200Hz.

0.3 Testing and Observations

Upon testing the voltmeter described in this report, it was seen that the voltmeter performed quite well. The RMS readings displayed on the board were always, without any known exception, accurate to within a very small margin from their expected values.

Due to the limitations on memory usage described in the [data processing section](#) above, the running window design forced display updates to occur rather slowly. For example, if the voltage on the ADC pin was changed, it would take approximately 10 seconds for the RMS voltage to stabilize on the display. Of course, this was to be expected considering the design that was implemented. However, due to this issue an alternate design likely would have made more sense - for example, if the RMS was taken over a 200ms interval and updated every 200ms.

Due to the FIR filtering and the fact that the running voltage is displayed as RMS, it was expected that the voltmeter would be quite accurate and would not show wild variations in the RMS, minimum, or maximum voltages once it has stabilized. In order to test the performance of the voltmeter, the following data was recorded when testing the voltmeter on various input voltages:

Table 3: Results from testing the voltmeter with various input voltages

Input voltage (V)	Stabilized RMS (V)	Minimum voltage (V)	Maximum voltage (V)	Percent error
0.00	0.00	0.00	0.00	0%
0.60	0.60	0.60	0.61	0%
0.60	0.61	0.60	0.61	1.64%
0.60	0.60	0.60	0.61	0%
1.00	1.01	1.01	1.01	1%
1.00	1.01	1.01	1.02	1%
1.00	1.01	1.00	1.02	1%
2.00	2.00	1.99	2.00	0%
2.00	2.01	2.00	2.01	0.5%
2.00	2.02	2.00	2.02	1%
2.56	2.57	2.56	2.57	0.39%
2.56	2.57	2.56	2.57	0.39%
2.56	2.57	2.56	2.58	0.39%
3.00	3.00	3.00	3.00	0%

It can be seen from [Table 3](#) that the voltmeter performed quite well. At 0V and 3V (the highest and lowest voltage pins on the board), the voltmeter performed perfectly. When the voltmeter was connected to the DAC for input for the other test points, there was some smaller error in some cases, but this error never even reaches 2%. As explained in the [data acquisition section](#), the resolutions of the ADC and the DAC were 8 bit, allowing for resolution down to approximately 0.012V. It is very possible that the error seen from the test data is due to a lack of precision due to the 8 bit resolution. This is further reinforced by the fact that the difference in voltage between the expected value and that output on the board's RMS display never exceeded 0.02V, regardless of the magnitude of the input voltage.

Overall, the voltmeter performed in complete accordance to the expectations. Although there was some small error in certain tests, the error was well within the range of the expected error due to the ADC and DAC resolution. Furthermore, as expected, the RMS voltage did converge 100% within 10 seconds, as the running window would suggest. Once the RMS voltage had stabilized, it remained very stable at its expected voltage (or close to the expected voltage, see [Table 3](#)), varying by at most 0.02V over a 30 second interval. The RMS did change slightly more than expected, even after convergence. This could have been due to insufficient filtering (perhaps 4th order was not enough after all). Finally, the minimum and maximum voltage readings remained exceptionally stable over the 30 second intervals that were tested for each test input. This was expected as particularly high or low voltage readings were likely filtered out by the FIR filter.

Appendix A

Circular Lists

In order to calculate RMS, minimum, and maximum on a moving window, a “circular” list data structure was used. Circular lists contain a capacity C describing how many items the list can hold, an index L describing where the next sample should go, and a vector δ which stores the data of the list. Each time a new element ϵ is to be added to the list, it is inserted by $\delta[(L++)\%C] = \epsilon$, where $\%$ represents the modulo operation, and $L++$ increments L . Note that when the circular list is full, that is to say $L = C$, the least recent data point in the list is overwritten. For example, if one was to add the values 1, 2, 3, 4, 5, 6 in a circular list of capacity 3, the progression of the circular list would look as follows:

$\delta = []$	Add 1
$\delta = [1,]$	Add 2
$\delta = [1, 2,]$	Add 3
$\delta = [1, 2, 3]$	Add 4
$\delta = [4, 2, 3]$	Add 5
$\delta = [4, 5, 3]$	Add 6
$\delta = [4, 5, 6]$	

Clearly, the δ in the previous example is always storing the latest 3 values that it receives. Therefore, the circular list is storing a running window of its latest C inputs.

This data structure is particularly convenient due to the running time of the operations it provides for storing running windows. To store a running window, one must add an item to a list and remove the oldest item from a list for each new sample in the worst case. With this data structure, adding a new item involves calculating an index by $(L++)\%C$, which is an $O(1)$ operation. However, adding the new item removes the oldest item implicitly! Therefore, the process of updating a running window with this data structure has $O(1)$ time complexity. In terms of storage, the circular list stores the C most recent samples as well as the value C itself and the current index into the list, L . Therefore, the space complexity of the circular list is $O(n)$, which cannot be improved. The addition of the two extra parameters L and C is a fairly low cost.