

ECSE 426 - Microprocessor Systems  
Lab Report 2: Timers, Interrupts, Multithreaded,  
Interrupt-Driven Readings and Peripheral Control

Harley Wiltzer (260690006)  
Matthew Lesko (260692352)

March 19, 2018

# Contents

1	Abstract . . . . .	1
2	Problem Statement . . . . .	1
3	Theory and Hypothesis . . . . .	2
3.1	Rectifier . . . . .	2
3.2	Multithreading and Semaphores . . . . .	2
3.3	Hypothesis . . . . .	2
4	Implementation . . . . .	3
4.1	PWM Pulse Generation . . . . .	3
4.2	ADC and Timers . . . . .	3
4.3	Filtering . . . . .	3
4.4	Alphanumeric Keypad . . . . .	4
4.5	Finite State Machine . . . . .	5
4.6	Voltage Controller . . . . .	7
4.7	Multithreading and System Architecture . . . . .	8
5	Testing and Observations . . . . .	9
5.1	PWM and Rectifier . . . . .	9
5.2	Filtering . . . . .	10
6	Conclusion . . . . .	11
<b>A GPIO Configuration Parameters</b>		<b>13</b>
<b>B ADC Configuration Settings</b>		<b>15</b>
<b>C TIM2 Configuration Settings</b>		<b>16</b>
<b>D TIM3 Configuration Settings</b>		<b>17</b>
<b>E HAL Cube MX Autogenerated Code</b>		<b>18</b>
<b>F HAL GPIO Pin Configuration</b>		<b>21</b>
<b>G Theory References</b>		<b>22</b>

# List of Figures

1	Finite State Machine . . . . .	5
2	System Architecture Block Diagram . . . . .	8
3	Filtered and Unfiltered Values VS Time (ms), Using Average FIR Filter . . . . .	11
4	Filtered and Unfiltered Values VS Time (ms), Using High Coefficients for Early Values	12
5	Filtered and Unfiltered Values VS Time (ms), Using High Coefficients for Late Values	13
F.1	HAL's Visualization of the STM32F4's Pins Configuration . . . . .	21

# List of Tables

1	Effect of input gestures on output . . . . .	6
2	Effect of rectifier parameters on its output voltage . . . . .	10

# 1 Abstract

The purpose of experiment 3 is for the programmers to gain experience in utilizing timers and interrupts to accomplish a task which involves converting an analog pulse to digital and displaying its voltage on an LED display, effectively a voltmeter. The purpose of experiment 4 is for the programmers to gain exposure in designing a multithreaded program on a real time operating system (RTOS) running on an embedded system. The task of experiment 4 involves copying over experiment 3's program and subdividing several of its features each to its own concurrently running thread, with the goal of optimizing power usage. This report will explain in detail how the programmers implemented the problems stated below, as well as the challenges they faced, the testing they had done, and the conclusions they have made. By the end of the report, the reader shall understand how the timers available on the STM32F4 board can be used to activate peripherals and generate a pulse, and understand the implementation of multithreaded programs on embedded systems.

## 2 Problem Statement

The problem is for the developers to first implement a solution for generating a PWM pulse whose duty cycle is controlled in order to maintain an approximately DC output voltage from a rectifier, where the voltage set by an input on a keypad. Afterwards, the device shall feed the rectifier's output to an ADC to be converted to a digital signal, and have the signal's mean voltage be automatically displayed on an LED display. Furthermore, the program has to be implemented with the use of concurrently-running threads running on an RTOS. The problem can be divided into the following tasks:

- Setting up a timer to act as a PWM pulse generator;
- Setting up a timer to activate the ADC to take an analog sample and convert it to digital;
- Designing a rectifier circuit component that takes the PWM pulse as input and feeds the output to the ADC;
- Testing and optimization of an FIR filter that reduces noise from the output signal of the ADC;
- Setting up the alphanumeric keypad so that the user may input their desired voltage to be displayed on an LED display;
- Maintaining the 7-segment display;
- Coding a controller function that automates the changes to be made on the PWM's duty cycle so that the current voltage adjusts to the target voltage set by the user;
- Coding a finite state machine function that enables the user to Enter and Delete digits for the target voltage, Reset the target voltage, and put the device to Sleep by using the keypad;
- Implementing the program's features using CMSIS-RTOS and multithreading;
- Reducing the power consumption of the device when it is in sleep mode using CMSIS-RTOS.

### 3 Theory and Hypothesis

Since this experiment uses components from experiment number 2, any theory for the identical components that has been mentioned in the previous lab report will be skipped. If you need to see theory for those components, and it is not mentioned here, please refer to Lab Report number 1.

#### 3.1 Rectifier

A rectifier is a component that converts alternating current (AC) to direct current (DC). It does so by only allowing a one-way flow of electrons, by the use of a diode. The diode allows electric current only in the forward bias condition and blocks electric current in reverse bias condition, this allows it to act like a rectifier. The output of the diode only contains a positive half cycle, as opposed to the input having a positive and a negative half cycle [1]. The rectifier component of this experiment's device contains a diode connected in series with a parallel connection of a resistor and a capacitor.

#### 3.2 Multithreading and Semaphores

The idea of multithreading is to have multiple threads execute concurrently. This allows for parallelism, more efficient use of a processor, and faster execution time. This is similar to the notion of concurrently-running processes, however, threads are not processes. One can have multiple threads running in one process, thus allowing for less overhead because the threads all share the same data [2]. Since, the threads may be sharing resources, this requires some kind of mutual exclusion to prevent threads from falling into deadlock or race-condition situations. The solution is semaphores. Semaphores allow one program to use a shared resource without having other programs use the resource at the same time. This is done by having a program wait until a semaphore is available to use before it can have access to the shared resource and execute its code. First a program waits for a semaphore, and once it is available, holds the semaphore by decrementing the semaphore's count, and executing its code. If the semaphore has a value of 0, no other program that is waiting for the same semaphore can execute. When the program using the shared resource no longer has need for the resource, it releases the semaphore by incrementing the semaphore's count, thus allowing other programs that are waiting to execute, to use the resource [2]. This resolves the problem of deadlocks and race-conditions.

#### 3.3 Hypothesis

The programmers should be able to implement a solution for experiment 3 given the time allocated for the project. The challenges one expects to face are learning how to use the keypad and implementing a solution with it, testing of the resistor and capacitor components for the rectifier, and implementing a solution for the different features available on the keypad with the use of a finite state machine. Furthermore, it is expected that a simple  $p$ -type controller design will be sufficient for adjusting the rectifier output voltage quickly. The programmers should be able to implement a solution for experiment 4 given the time allocated for the project, since the majority of the device and code is the same from experiment 3 and the students should have had experience with multithreading and semaphores. Due to the independence of various components of the system, it is expected for the code from experiment 3 to be easily divisible into several concurrent

threads, which would increase the efficiency of the system and readability of the code without any non-negligible performance sacrifice.

## 4 Implementation

### 4.1 PWM Pulse Generation

The developers were tasked with designing a system that generates PWM pulses and feeds the voltage to a rectifier circuit. Under the [TIM3 configuration settings](#), a timer, one can see that it is configured to generate a pulse at a frequency of 500kHz, by inheriting the APB clock of 84MHz and having a period of 168, hence having 84MHz divided by 168, yielding 500kHz. The students have chosen to use this frequency after having tested lower frequencies, and came to the conclusion that in order for the duty cycle to control the output voltage, the frequency had to be elevated to at least this amount of frequency. After [testing and observing](#), the students noticed that a frequency of 500kHz was suitable enough to have the duty cycle control the output voltage. The STM32F4's TIM3 hardware timer is configured to generate and output a PWM pulse channel that is assigned to a GPIO pin on the board, which can be viewed in the [GPIO pin configuration table](#) generated by STMCubeMX. By using a circuit wire, the programmers can feed the output of the PWM pulse channel to a circuit component on a bread board, more specifically, the rectifier.

The rectifier circuit holds the charge in the capacitor while a load resistance discharges the capacitor, enabling the user to influence the duty cycle to control the output voltage level. The required range for the output voltage is between 0.5V and 2.8V, hence it is required to test different configurations for the resistor and capacitor to deliver a range of output voltages that meet the requirement. The students have chosen a capacitor of 5uF and a resistor of 390 Ohms. This configuration was chosen because it achieved an output voltage range of 0.4V to roughly 2.1V. The programmers were able to observe this phenomenon after [testing](#) multiple different configurations and came to the conclusion that the resistor and capacitor values chosen were suitable for the task.

### 4.2 ADC and Timers

Previously, the students used STM's built-in SysTick to enable the ADC. For this experiment, the students had configured the ADC to be triggered by an STM32F4 timer, specifically TIM2. Under the [TIM2 configuration settings](#), one can see that the timer is configured with a prescaler of 83999 and a period of 1, which by inheriting the APB clock of 84MHz, entails the timer to a frequency of 1kHz. The reason behind this decision is so that the ADC gets activated frequently enough to take samples more often. This would allow it to be a more responsive component and allow the output voltage's RMS to be updated more accurately. One can see under the [ADC's configuration settings](#) that the trigger for the ADC's activation is the TIM2's trigger event.

### 4.3 Filtering

The device uses a modified FIR filter from experiment 2 in order to reduce the noise of the ADC's signal. The two modifications are the following:

- Increasing from 5 coefficients to 10 being used in the moving average;
- Setting the first five coefficients to 0.05 and the last five coefficients to 0.15.

In this way, the 5 earliest samples have a significance of 0.05 and the 5 later samples have a significance of 0.15 when calculating the average of the 10 samples. The programmers drew this conclusion by trial and error and the empirical evidence proved that the above modifications to the FIR filter significantly reduced the signal's noise. The programmers have tested multiple configurations, in which this report shall demonstrate three of the tested configurations. The graphs the programmers have used to determine the optimized filter can be seen in the [Testing and Observations: Filtering](#), section. After testing multiple configurations, the three that are graphed in the testing section further prove that the programmers have chosen an accurate filter with coefficients: [0.15, 0.15, 0.15, 0.15, 0.15, 0.05, 0.05, 0.05, 0.05].

## 4.4 Alphanumeric Keypad

In order to allow a user to either control the output voltage of the rectifier circuit or put the system to sleep, it was required to integrate a keypad peripheral. The keypad has buttons arranged in a  $4 \times 3$  matrix which is interpreted via the 7 pins that the keypad offers. In order to read a given column of the keypad, one may set the keypad's corresponding column pin HIGH and subsequently read all four row pins. The reading of the four row pins and the column pins are passed through a lookup table (implemented in software) which returns the corresponding key that was pressed (assuming only one key is pressed at any given time).

The overall strategy for processing keyboard input is as follows:

1. Set the pin corresponding to the first keypad column HIGH.
2. Poll the row pins for a short period (on the order of 100ms).
3. If the row pins indicate that a button was pressed, search lookup table for the appropriate button reading and latch it in a `reading` variable
4. Repeat steps 2 and 3 for each of the remaining column pins.
5. Once all column pins have been read, alert other tasks of the `reading` variable (via a signal in the context of an RTOS, for example) and reset `reading = -1`. Note that it is important to alert other tasks when no button has been pressed as well (indicated by `reading = -1`) in order to distinguish between a user pressing a button and a user holding a button.
6. Repeat

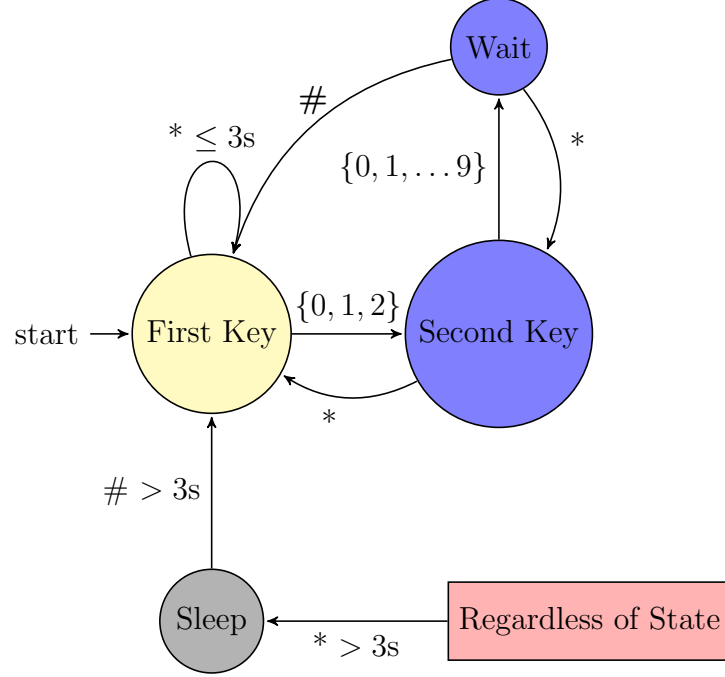
The above procedure assumes the keypad controller is operating concurrently with the rest of the system, which was true in the case of Lab 4 which made use of an RTOS. However, during lab 3, this was not at all the case. Since no signals could be sent, the finite state machine would wait for the whole keyboard to be scanned before computing its next state. Clearly, this is rather inefficient as the finite state machine would, first of all, have to wait approximately 300ms before each state transition. Secondly, since most of the time the user is not pressing a button, the finite state machine would end up polling for button presses needlessly for the majority of the CPU time. With the threading and signal implements in FreeRTOS, this efficiency was increased greatly. In the implementation of Lab 4, as discussed below.

## 4.5 Finite State Machine

A finite state machine was implemented to control the functionality of the system according to the state the system is in. The state is changed by inputs to the keypad from the user.

Figure 1 shows the semantics of the state transitions, which was kept exactly the same be-

Figure 1: Finite State Machine



*Yellow states: display shows measured voltage reading*  
*Blue states: display shows current voltage typed by user*  
*Gray states: display is off*

tween labs 3 and 4. The arrows with labels that have inequalities are triggered when the symbol before the inequality is held for a duration that satisfies the inequality. For example, the arrow from “Sleep” to “First Key” with the label “ $\# > 3s$ ” should be interpreted as “when in the Sleep state, if the  $\#$  key is held for more than 3 seconds, transition to the First Key state”.

Although the FSM shown shows how the states are changed, details about the output of the system are omitted in the diagram to avoid over-cluttering it. The details about how the output of the system changes are listed in Table 1. Although the state transitions were implemented identically in Labs 3 and 4, the RTOS used in Lab 4 allowed for a much more efficient FSM procedure, as was foreshadowed above. Since the FSM must wait for a new keyboard reading before transitioning states, in Lab 3 this involved *busy waiting* before control was passed to the FSM. The SysTick timer was used to time the keyboard reading, so the FSM would essentially wait for SysTick to count up to the value corresponding to the end of the third column of the keypad’s reading period before processing any potential state transitions. This wastes CPU time and energy.

Using multithreading capabilities of FreeRTOS, this problem was solved in Lab 4. The keypad processing task and FSM control task were each given their own thread, so they could run



Table 1: Effect of input gestures on output

Source State	Action	Destination State	Output
Wait	Press #	First Key	Controller adjusts output voltage to that specified by the user
Wait	Press *	Second Key	Second input digit is erased
Second Key	Press *	First Key	First input digit is erased
<i>Any</i>	Hold * for around 1s	First Key	Output voltage is set to 0V
<i>Any</i>	Hold * for > 3s	Sleep	Peripherals turned off, low energy mode
Sleep	Hold # for > 3s	First Key	Peripherals and threads restarted

concurrently. Furthermore, semaphores and signals were used to avoid busy waiting. Firstly, to avoid busy waiting at the FSM, the FSM would always wait on a *signal* from the keypad before proceeding. This puts the FSM in the Waiting state, where it will not be scheduled by the kernel. This frees up CPU time for other tasks. Whenever the keypad task process a key reading (whether that be an actual button press, or a detection that no button is being held), the keypad task signals the FSM task, putting its thread back in the Ready state, where it can once again be scheduled by the kernel. Immediately after receiving the signal, the FSM sets its signal back to 0, to force itself to wait for another signal from the keypad before the next iteration.

Although this multithreading idea does increase the efficiency of certain parts of the system, it brings along some undesirable consequences as well. The problem of mutual exclusion is of particular significance: if two threads run concurrently and certain data shared between the two will be written to by the threads, the value of the data will be indeterministic and will depend on the timing of when the threads are scheduled. Thus, we must provide a way to prevent both threads from accessing these resources at the same time when one of them is writing to it. Additionally, this mediation must not allow the processes to ever both wait for each other resulting in an “infinite loop”, called deadlock. A common solution to this problem is the use of semaphores. Semaphores normally contain an integer  $i$  that counts how many free units of the resources that it controls are available and a queue  $q$  that stores all threads that are waiting on resources to be free. When a thread “waits” on a semaphore, we have  $i \leftarrow i - 1$ , and the proceeding behavior then depends on the sign of  $i$ :

- When  $i \geq 0$ , the thread continues as usual
- When  $i < 0$ , the thread enters the Waiting state (so it can no longer be scheduled by the dispatcher) and is enqueued in  $q$

Clearly, as in the case of the signal, this efficiently allows a thread to wait until a resource is available. When a thread is finished with a shared resource, it “releases” the semaphore, which causes  $i \leftarrow i + 1$ . Again, the proceeding behavior depends on  $i$ :

- When  $i = 1$ ,  $x \leftarrow \text{dequeue}(q)$ , and put  $x$  in the Ready state to allow it to be scheduled once again

- Otherwise, continue like usual

This shows how threads are given permission to access a shared resource. Since this method prevents busy waiting like in the case of the signal, it is a fairly efficient way to ensure that only one thread can access a shared resource at a time.

In the system of Lab 4, the FSM shared one resource with the keypad task: the keypad reading itself. To protect this shared resource, a binary semaphore was used, as the keypad reading is one resource and thus it is either free or not free. When one of the FSM thread or the keypad thread are accessing the keypad reading data, the other thread is put in the Waiting state to avoid wasting CPU time. Additionally, race conditions are avoided, and there is no corruption in the key reading data as seen by either thread. To allow both the FSM and keypad threads to run concurrently, the FSM latches the keypad reading and stores it in another unshared variable while it has access to the keypad reading, to effectively allow the keypad thread to continue toying with the keypad reading while the FSM executes its tasks. This was especially convenient given the vastly different frequencies between the two threads.

## 4.6 Voltage Controller

The main requirement of the system in Labs 3 and 4 was to maintain a reasonably constant voltage output, independent of the load on the circuit, and allow this constant voltage to be determined by the user in real time. To accomplish this the duty cycle of the PWM output of the board's TIM3 PWM timer was controlled and sent through a rectifier to provide the desired load-independent voltage.

The TIM3 PWM timer counts up to a given period value  $T$ , at a counting rate determined by a given prescaler  $P$ . As discussed above,  $T$  and  $P$  were chosen such that the period of the PWM pulse was 500kHz. Furthermore, the TIM3 PWM timer has another parameter,  $\delta$ , which controls the duty cycle of the PWM pulse. Whenever the count of the timer exceeds or falls below  $\delta$ , the PWM output toggles. Therefore, the fraction of the period for which the PWM pulse is high (in other words, the duty cycle) is determined by the ratio  $\delta : T$ .

When  $\delta$  is large, the rectifier circuit has more time to charge its capacitor, and therefore its output voltage increases. On the other hand, when  $\delta$  is low, the rectifier spends most of its time discharging the capacitor and decreasing the voltage across it, yielding a lower output voltage.

In order to set  $\delta$  appropriately to provide the desired rectifier output, feedback from the ADC is used. Therefore, action is taken in the ADC's conversion callback function, which is signaled in the interrupt triggered by an ADC conversion. To modify  $\delta$ , a  $p$ -type controller was implemented. We have, after every ADC conversion,

$$\delta \leftarrow \delta + p(t - x) \tag{1}$$

Where  $t$  is the target voltage,  $x$  is the latest (filtered) voltage reading, and  $p$  is a proportionality constant chosen experimentally to be 0.01. The main criteria for choosing  $p$  were the convergence time of the controller and its stability. For lower values of  $p$ , it was found that the controller took longer to converge to an appropriate output voltage. However, for greater values of  $p$ , the controller would oscillate about the desired voltage too much, resulting in a less constant output. The value

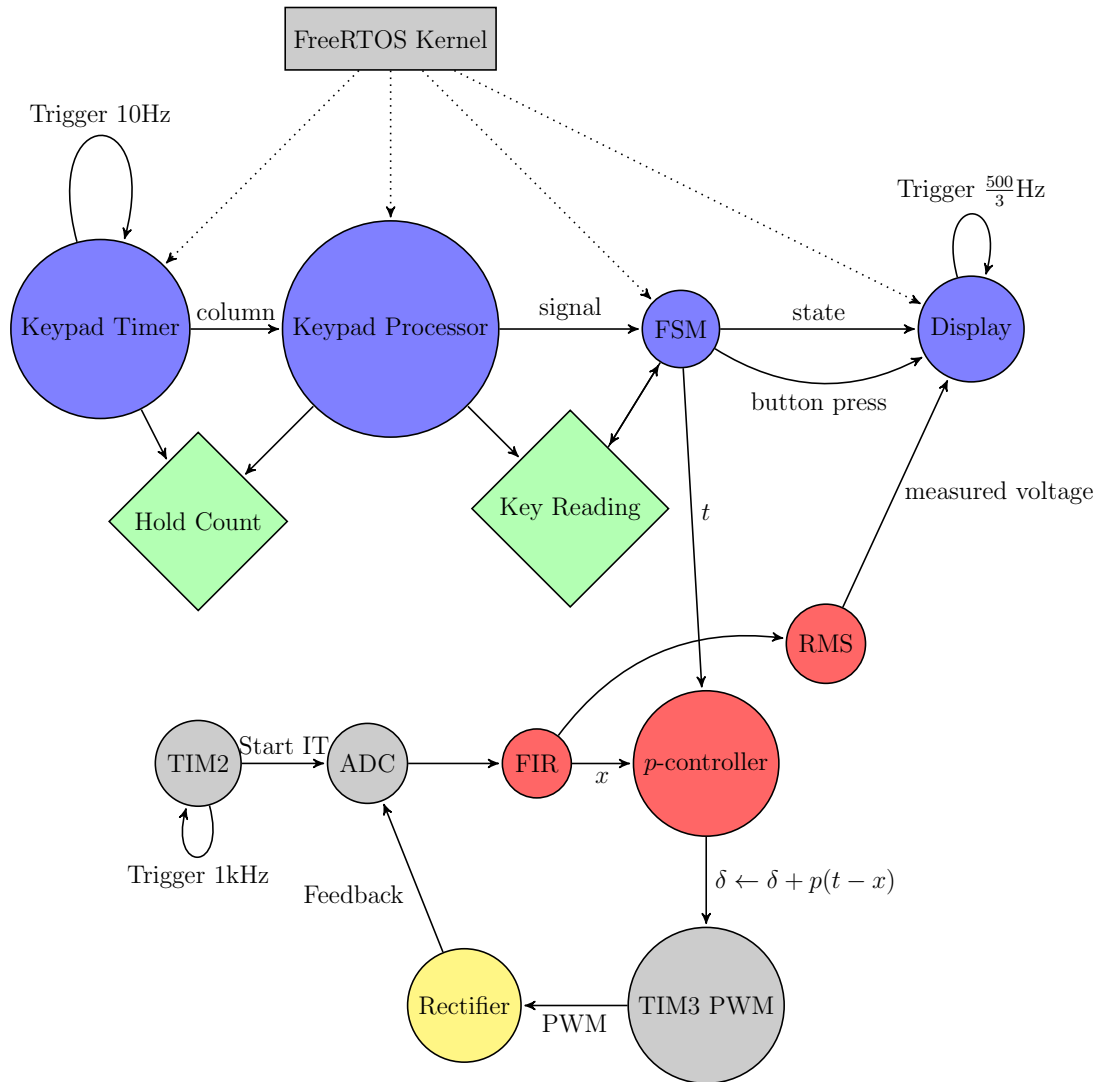
$p = 0.01$  was found to be a good balance between speed and stability. During the testing phase, it was found that the system converged to the appropriate voltage reading in approximately 40ms on average.

## 4.7 Multithreading and System Architecture

Figure 2 shows the overall architecture of the system, including the organization of the threads in Lab 4.

The design of the threads of for the FSM and keypad tasks have been discussed above. However,

Figure 2: System Architecture Block Diagram



*Blue circles are threads*  
*Green rhombi are shared resources*  
*Red circles are software implementations*  
*Yellow circles are electric circuits*

certain designs of Lab 4 led to some interesting thread usage for the display and keypad timer tasks. Firstly, powering the 7 segment display was very inefficient in Lab 3. In order to control the timing of the refresh rates of the individual displays, display logic was included in the SysTick interrupt callback. This is bad practice, as the SysTick callback should be as “lightweight” as possible to ensure the desired frequency of the interrupt. However, given the opportunity to use threads, the display logic was implemented in its own thread entirely by OS timers to diminish busy waiting. This once again frees up CPU time for other tasks and does not interfere with the interrupt routine of any other time-sensitive interrupt.

Another benefit of using threads is that it makes the system more modular, which allowed more control of which components were consuming power in the system. Therefore, an attempt was made to decrease power usage when the device is in sleep mode. In sleep mode, the FSM and the display are irrelevant, so their respective threads were stopped - preventing them from drawing power. Furthermore, the peripherals such as the ADC and the timers were stopped to decrease power usage in the dormant state. However, a more effective method would have been to deinitialize the peripherals, which has a much greater effect on power usage.

However, with the FSM thread off in sleep mode, logic for determining when the # key was held for 3 seconds to wake up the system had to be taken care of somewhere else, as it was initially implemented in the FSM for Lab 3. Luckily, it was fairly simple to move this logic to the keypad timer thread. This modification required another binary semaphore to protect a wakeup counter that was shared between the keypad timer and the keypad processor tasks. Note that the keypad and keypad timer threads must remain on while the device is in sleep mode, because without them the device cannot be woken up!

Of course, the efforts made to reduce power consumption must be reversed when the device is woken up. When the keypad timer thread detects a wakeup, it resets the state to First Key, and re-creates the thread ID's that were defined globally. Furthermore, the peripherals such as the ADC and the timers are started once again.

## 5 Testing and Observations

### 5.1 PWM and Rectifier

The students tested multiple configurations of resistor and capacitor values for the rectifier circuit in order to determine which parameters allowed the largest range in DC output voltage. Resistances in the range of  $390\Omega - 4.7k\Omega$  and capacitances in the range of  $5\mu F - 20\mu F$  were tested. These ranges were chosen since the time constants generated by those parameters satisfy  $(RC)^{-1} \ll f$ , where  $f$  is the PWM frequency. The time constant is a measure that is related to how long the circuit takes to charge and discharge. It was desired for this length to be far greater than the PWM frequency so that the circuit would not have enough time to charge and discharge by vast amounts, thereby keeping its output voltage fairly constant.

One can see from the results shown in [Table 2](#) that having a small resistor (390 Ohms) and low capacitor (5uF) allow the output voltage to yield a range roughly between 0.4V and 2.1V,

Table 2: Effect of rectifier parameters on its output voltage

Resistance (k $\Omega$ )	Capacitance( $\mu$ F)	Duty Cycle (%)	Output Mean Voltage (V)
4700	20	0.05	1.620
4700	20	0.95	2.150
4700	5	0.05	1.600
4700	5	0.95	2.140
390	20	0.05	0.364
390	20	0.95	1.920
390	20	1.00	1.920
390	5	0.05	0.401
390	5	0.95	1.930
390	5	1.00	2.105

making it the best suitor for the requirement of delivering a wide range of output voltages. This configuration also has another important advantage. Say a load resistance  $\tilde{R}$  is added to the circuit. In this case, the equivalent resistance connected across the capacitor of the rectifier circuit is

$$R||\tilde{R} \equiv \frac{R\tilde{R}}{R + \tilde{R}} \leq \min(R, \tilde{R}) \quad (2)$$

where  $R$  is the resistor used in the rectifier circuit. Say we wish that this equivalent resistance is greater than  $\alpha R$  (where  $\alpha < 1$ ), that is to say, adding the load resistance doesn't affect the resistance of the circuit by more than some constant threshold. Then,

$$\frac{R||\tilde{R}}{R} = \frac{\tilde{R}}{R + \tilde{R}} \geq \alpha \rightarrow \left(\frac{1 - \alpha}{\alpha}\right) \tilde{R} \geq R \quad (3)$$

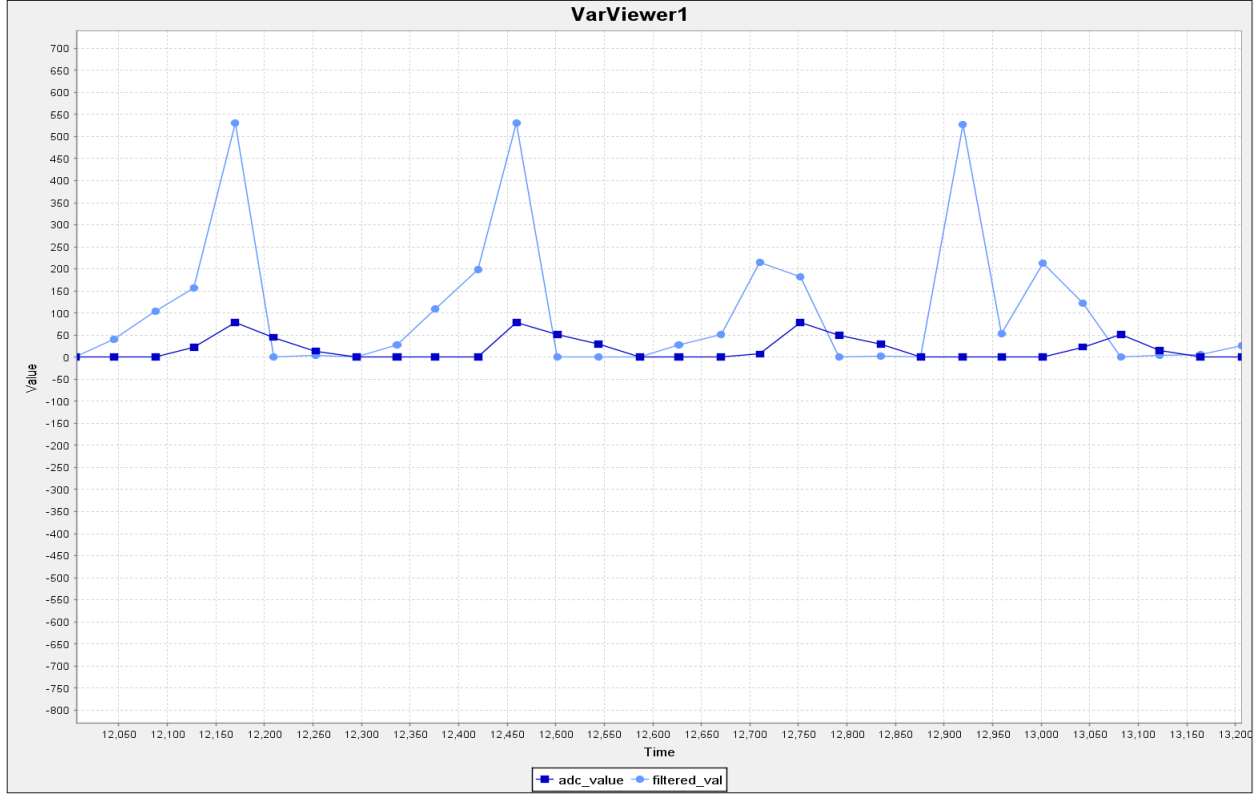
Equation (3) shows that as  $R$  grows, the smallest  $\tilde{R}$  that can be loaded to the circuit before changing the equivalent resistance by the  $\alpha$  threshold grows linearly. Therefore, by choosing a small  $R$ , the range of load resistances that have little effect on the equivalent circuit is greater! This means that for a larger set of load resistances, the loading on the circuit will have very little effect on the circuit itself, so it can be expected that the rectifier circuit's behavior will not vary greatly by the load resistance. Therefore, a small  $R$  for the rectifier was a good choice for two main reasons: it provided a good range of output voltages, and it is less sensitive to loading effects than a larger resistance.

## 5.2 Filtering

One can draw comparisons from the following graphs generated by a run-time variables monitoring and visualization tool, STM Studio[1]:

In Figure 3, the configuration for the filter's coefficients is [0.2, 0.2, 0.2, 0.2, 0.2]. In regular blue are the unfiltered values, and in light blue are the filtered values. This graph demonstrates the values read at runtime of the unfiltered and filtered values that passed through the unmodified FIR filter. One can observe that there is a considerable amount of noise left from filtering.

Figure 3: Filtered and Unfiltered Values VS Time (ms), Using Average FIR Filter



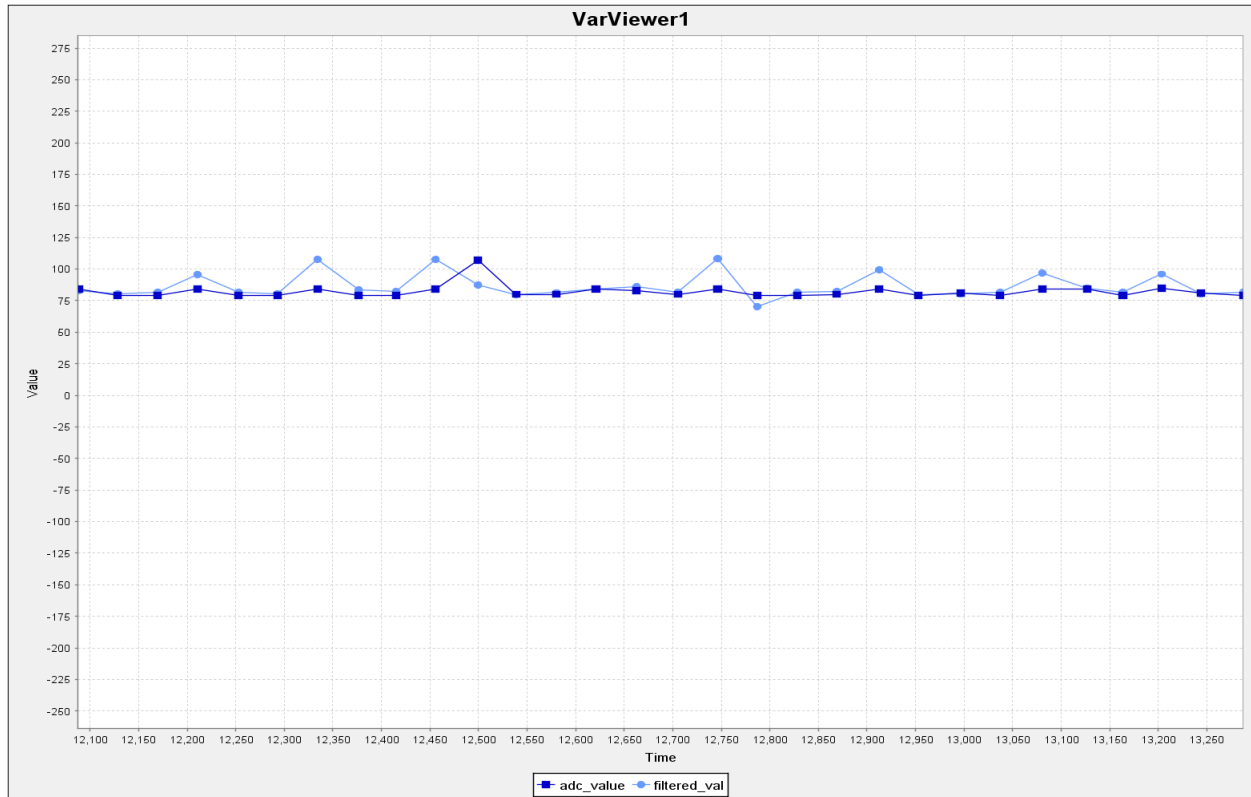
In Figure 4, the configuration for the filter's coefficients is  $[0.05, 0.05, 0.05, 0.05, 0.05, 0.15, 0.15, 0.15, 0.15, 0.15]$ . This filter considers the five earliest values to each have a significance of 0.15 and the five later values to each have a significance of 0.05. One can observe by the graph that this filter is a considerable improvement to the unmodified version. Although, the programmers believe that it could be optimized further.

Finally, in Figure 5, the configuration for the filter's coefficients is  $[0.15, 0.15, 0.15, 0.15, 0.15, 0.05, 0.05, 0.05, 0.05, 0.05]$ . This filter considers the five later values to each have a significance of 0.15 and the five early values to each have a significance of 0.05. One can observe by the graph that this filter is an improvement to the previous version. The programmers believe that this version of the filter should perform well enough given the scope of the problem.

## 6 Conclusion

It can be concluded that the STM32 board along with a simple rectifier circuit are capable of creating a configurable DC voltage source with considerably high resistance to loading effects. By tuning the frequency of the PWM output of TIM3, it was possible to generate a reasonably flat output voltage from a square wave PWM input signal. Furthermore, adjusting TIM3's `Pulse` parameter (and thereby its duty cycle, as discussed previously) gave a sufficient range of approximately DC output voltages from the rectifier. Interestingly enough, a *very* simple control system, the *p*-type controller, was able to automatically adjust the duty cycle extremely quickly, and the

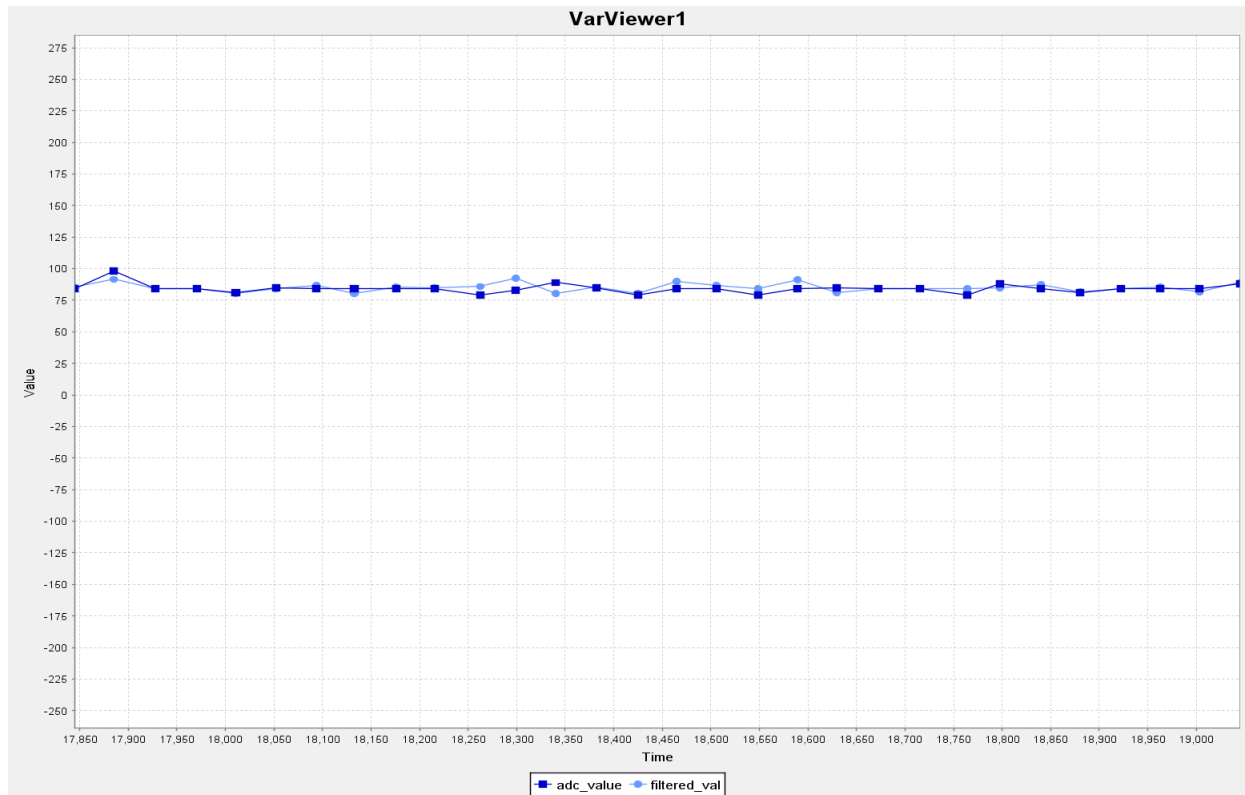
Figure 4: Filtered and Unfiltered Values VS Time (ms), Using High Coefficients for Early Values



rate at which the output voltage converged was quite good.

After implementing this system with and without a kernel and operating system, it was immediately clear that the RTOS was very handy for improving the efficiency of the system. Certain components, particularly the display and the keyboard, operate independently of most other components and allowing them to run concurrently at their own frequencies was not only more CPU-efficient, but was simpler to implement as well. Using several threads allowed a more modular design, making the code actually easier to read and reason about in Lab 4 than Lab 3. Furthermore, the capacity to prevent certain components from being scheduled by the CPU was a nice way to improve the efficiency of the system, especially in the case of the FSM which always needs to wait for a keypad reading before processing. Since a complete scan of the keypad is far slower than a state transition, most of the CPU's time was wasted in the Lab 3 implementation, but using signals eliminated this issue in Lab 4. Finally, it was possible to easily turn off threads, eliminating their power consumption. Although the great power of threads did indeed come with great responsibility (mutual exclusion, for example), the experience of implementing it was instructive and overall led to a cleaner and more intuitive system design.

Figure 5: Filtered and Unfiltered Values VS Time (ms), Using High Coefficients for Late Values



## Appendix A

### GPIO Configuration Parameters

This appendix lists the configuration parameters set for each of the different GPIO pins (or classes of GPIO pins).

#### User Input Button

Parameter	Value
Mode	GPIO_MODE_IT_RISING
Pull	GPIO_NOPULL



**Display Mode LEDs (4 of these)**

Parameter	Value
Mode	GPIO_MODE_OUTPUT_PP
Pull	GPIO_NOPULL
Speed	GPIO_SPEED_FREQ_LOW

**Display Segment Pins (8 of these)**

Parameter	Value
Mode	GPIO_MODE_OUTPUT_PP
Pull	GPIO_NOPULL
Speed	GPIO_SPEED_FREQ_LOW

**Display Selector Pins (3 of these)**

Parameter	Value
Mode	GPIO_MODE_OUTPUT_PP
Pull	GPIO_NOPULL
Speed	GPIO_SPEED_FREQ_LOW

# Appendix B

## ADC Configuration Settings

### ADC Instance Parameters

Parameter	Value
Clock Prescaler	ADC_CLOCK_SYNC_PCLK_DIV4
Resolution	ADC_RESOLUTION_8B
Scan Conversion Mode	Disabled
Continuous Conversion Mode	Disabled
Discontinuous Conversion Mode	Disabled
External Trigger Conversion Edge	ADC_EXTERNALTRIGCONVEDGE_RISING
External Trigger Conversion	ADC_EXTERNALTRIGCONV_T2_TRGO
Data Alignment	ADC_DATAALIGN_RIGHT
Number of Conversions	1
DMA Continuous Requests	Disabled
EOC Selection	ADC_EOC_SINGLE_CONV

### ADC Channel Parameters (Channel 1)

Parameter	Value
Rank	1
Sampling Time	ADC_SAMPLETIME_28CYCLES

# Appendix C

## TIM2 Configuration Settings

### TIM2 Instance Parameters

Parameter	Value
Instance	TIM2
Clock Prescaler	83999
Counter Mode	TIM_COUNTERMODE_UP
Period	1
Clock Division	TIM_CLOCKDIVISION_DIV1

### TIM2 Clock Source Parameters

Parameter	Value
Clock Source	TIM_CLOCKSOURCE_INTERNAL

### TIM2 Master Configuration Parameters

Parameter	Value
Master Output Trigger	TIM_TRGO_UPDATE
Master Slave Mode	TIM_MASTERSLAVEMODE_DISABLE

# Appendix D

## TIM3 Configuration Settings

### TIM3 Instance Parameters

Parameter	Value
Instance	TIM3
Clock Prescaler	0
Counter Mode	TIM_COUNTERMODE_UP
Period	PWM_PERIOD
Clock Division	TIM_CLOCKDIVISION_DIV1

### TIM3 Master Configuration Parameters

Parameter	Value
Master Output Trigger	TIM_TRGO_RESET
Master Slave Mode	TIM_MASTERSLAVEMODE_DISABLE

### TIM3 Output Channel Parameters

Parameter	Value
OC Mode	TIM_OCMODE_PWM1
Pulse	duty_cycle * PWM_PERIOD
OC Polarity	TIM_OCPOLARITY_HIGH
OC Fast Mode	TIM_OCFAST_DISABLE

# Appendix E

## HAL Cube MX Autogenerated Code

```
GPIO_InitTypeDef GPIO_InitStructure;

/* GPIO Ports Clock Enable */
__HAL_RCC_GPIOE_CLK_ENABLE();
__HAL_RCC_GPIOC_CLK_ENABLE();
__HAL_RCC_GPIOH_CLK_ENABLE();
__HAL_RCC_GPIOA_CLK_ENABLE();
__HAL_RCC_GPIOB_CLK_ENABLE();
__HAL_RCC_GPIOD_CLK_ENABLE();

/*Configure GPIO pin Output Level */
HAL_GPIO_WritePin(CS_I2C_SPI_GPIO_Port, CS_I2C_SPI_Pin, GPIO_PIN_RESET);

/*Configure GPIO pin Output Level */
HAL_GPIO_WritePin(OTG_FS_PowerSwitchOn_GPIO_Port, OTG_FS_PowerSwitchOn_Pin, GPIO_PIN_SET);

/*Configure GPIO pin Output Level */
HAL_GPIO_WritePin(GPIOD, LD4_Pin | LD3_Pin | LD5_Pin | LD6_Pin | Audio_RST_Pin, GPIO_PIN_RESET);

/*Configure GPIO pin : CS_I2C_SPI_Pin */
GPIO_InitStructure.Pin = CS_I2C_SPI_Pin;
GPIO_InitStructure.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStructure.Pull = GPIO_NOPULL;
GPIO_InitStructure.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(CS_I2C_SPI_GPIO_Port, &GPIO_InitStructure);

/*Configure GPIO pin : OTG_FS_PowerSwitchOn_Pin */
GPIO_InitStructure.Pin = OTG_FS_PowerSwitchOn_Pin;
GPIO_InitStructure.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStructure.Pull = GPIO_NOPULL;
GPIO_InitStructure.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(OTG_FS_PowerSwitchOn_GPIO_Port, &GPIO_InitStructure);

/*Configure GPIO pin : PDM_OUT_Pin */
GPIO_InitStructure.Pin = PDM_OUT_Pin;
GPIO_InitStructure.Mode = GPIO_MODE_AF_PP;
GPIO_InitStructure.Pull = GPIO_NOPULL;
GPIO_InitStructure.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStructure.Alternate = GPIO_AF5_SPI2;
HAL_GPIO_Init(PDM_OUT_GPIO_Port, &GPIO_InitStructure);

/*Configure GPIO pin : B1_Pin */
GPIO_InitStructure.Pin = B1_Pin;
GPIO_InitStructure.Mode = GPIO_MODE_IT_RISING;
GPIO_InitStructure.Pull = GPIO_NOPULL;
HAL_GPIO_Init(B1_GPIO_Port, &GPIO_InitStructure);

/*Configure GPIO pin : I2S3_WS_Pin */
GPIO_InitStructure.Pin = I2S3_WS_Pin;
GPIO_InitStructure.Mode = GPIO_MODE_AF_PP;
GPIO_InitStructure.Pull = GPIO_NOPULL;
GPIO_InitStructure.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStructure.Alternate = GPIO_AF6_SPI3;
```

```

HAL_GPIO_Init(I2S3_WS_GPIO_Port, &GPIO_InitStruct);

/*Configure GPIO pins : SPI1_SCK_Pin SPI1_MISO_Pin SPI1_MOSI_Pin */
GPIO_InitStruct.Pin = SPI1_SCK_Pin | SPI1_MISO_Pin | SPI1_MOSI_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_AF_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStruct.Alternate = GPIO_AF5_SPI1;
HAL_GPIO_Init(GPIOA, &GPIO_InitStruct);

/*Configure GPIO pin : BOOT1_Pin */
GPIO_InitStruct.Pin = BOOT1_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_INPUT;
GPIO_InitStruct.Pull = GPIO_NOPULL;
HAL_GPIO_Init(BOOT1_GPIO_Port, &GPIO_InitStruct);

/*Configure GPIO pin : CLK_IN_Pin */
GPIO_InitStruct.Pin = CLK_IN_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_AF_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStruct.Alternate = GPIO_AF5_SPI2;
HAL_GPIO_Init(CLK_IN_GPIO_Port, &GPIO_InitStruct);

/*Configure GPIO pins : LD4_Pin LD3_Pin LD5_Pin LD6_Pin
Audio_RST_Pin */
GPIO_InitStruct.Pin = LD4_Pin | LD3_Pin | LD5_Pin | LD6_Pin | Audio_RST_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(GPIOD, &GPIO_InitStruct);

/*Configure GPIO pins : I2S3_MCK_Pin I2S3_SCK_Pin I2S3_SD_Pin */
GPIO_InitStruct.Pin = I2S3_MCK_Pin | I2S3_SCK_Pin | I2S3_SD_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_AF_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStruct.Alternate = GPIO_AF6_SPI3;
HAL_GPIO_Init(GPIOC, &GPIO_InitStruct);

/*Configure GPIO pin : VBUS_FS_Pin */
GPIO_InitStruct.Pin = VBUS_FS_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_INPUT;
GPIO_InitStruct.Pull = GPIO_NOPULL;
HAL_GPIO_Init(VBUS_FS_GPIO_Port, &GPIO_InitStruct);

/*Configure GPIO pins : OTG_FS_ID_Pin OTG_FS_DM_Pin OTG_FS_DP_Pin */
GPIO_InitStruct.Pin = OTG_FS_ID_Pin | OTG_FS_DM_Pin | OTG_FS_DP_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_AF_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStruct.Alternate = GPIO_AF10_OTG_FS;
HAL_GPIO_Init(GPIOA, &GPIO_InitStruct);

/*Configure GPIO pin : OTG_FS_OverCurrent_Pin */
GPIO_InitStruct.Pin = OTG_FS_OverCurrent_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_INPUT;
GPIO_InitStruct.Pull = GPIO_NOPULL;
HAL_GPIO_Init(OTG_FS_OverCurrent_GPIO_Port, &GPIO_InitStruct);

/*Configure GPIO pins : Audio_SCL_Pin Audio_SDA_Pin */
GPIO_InitStruct.Pin = Audio_SCL_Pin | Audio_SDA_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_AF_OD;
GPIO_InitStruct.Pull = GPIO_PULLUP;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
GPIO_InitStruct.Alternate = GPIO_AF4_I2C1;
HAL_GPIO_Init(GPIOB, &GPIO_InitStruct);

/*Configure GPIO pin : MEMS_INT2_Pin */
GPIO_InitStruct.Pin = MEMS_INT2_Pin;
GPIO_InitStruct.Mode = GPIO_MODE_EVT_RISING;

```

```

GPIO_InitStruct.Pull = GPIO_NOPULL;
HAL_GPIO_Init(MEMS_INT2_GPIO_Port, &GPIO_InitStruct);

/*Configure 7 segment pins*/
GPIO_InitStruct.Pin = GPIO_PIN_7 | GPIO_PIN_11 | GPIO_PIN_10 | GPIO_PIN_14 | GPIO_PIN_15;
GPIO_InitStruct.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(GPIOE, &GPIO_InitStruct);
GPIO_InitStruct.Pin = GPIO_PIN_12 | GPIO_PIN_13;
GPIO_InitStruct.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(GPIOB, &GPIO_InitStruct);
GPIO_InitStruct.Pin = GPIO_PIN_8;
GPIO_InitStruct.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(GPIOD, &GPIO_InitStruct);

/*Configure digit selector pins*/
GPIO_InitStruct.Pin = GPIO_PIN_2 | GPIO_PIN_3 | GPIO_PIN_0;
GPIO_InitStruct.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(GPIOD, &GPIO_InitStruct);

/*Configure keypad column pins*/
GPIO_InitStruct.Pin = GPIO_PIN_5 | GPIO_PIN_3 | GPIO_PIN_1;
GPIO_InitStruct.Mode = GPIO_MODE_OUTPUT_PP;
GPIO_InitStruct.Pull = GPIO_NOPULL;
GPIO_InitStruct.Speed = GPIO_SPEED_FREQ_LOW;
HAL_GPIO_Init(GPIOE, &GPIO_InitStruct);

/*Configure keypad row pins*/
GPIO_InitStruct.Pin = GPIO_PIN_8 | GPIO_PIN_7 | GPIO_PIN_5;
GPIO_InitStruct.Mode = GPIO_MODE_INPUT;
GPIO_InitStruct.Pull = GPIO_PULLDOWN;
HAL_GPIO_Init(GPIOB, &GPIO_InitStruct);
GPIO_InitStruct.Pin = GPIO_PIN_7;
GPIO_InitStruct.Mode = GPIO_MODE_INPUT;
GPIO_InitStruct.Pull = GPIO_PULLDOWN;
HAL_GPIO_Init(GPIOD, &GPIO_InitStruct);

```

# Appendix F

## HAL GPIO Pin Configuration

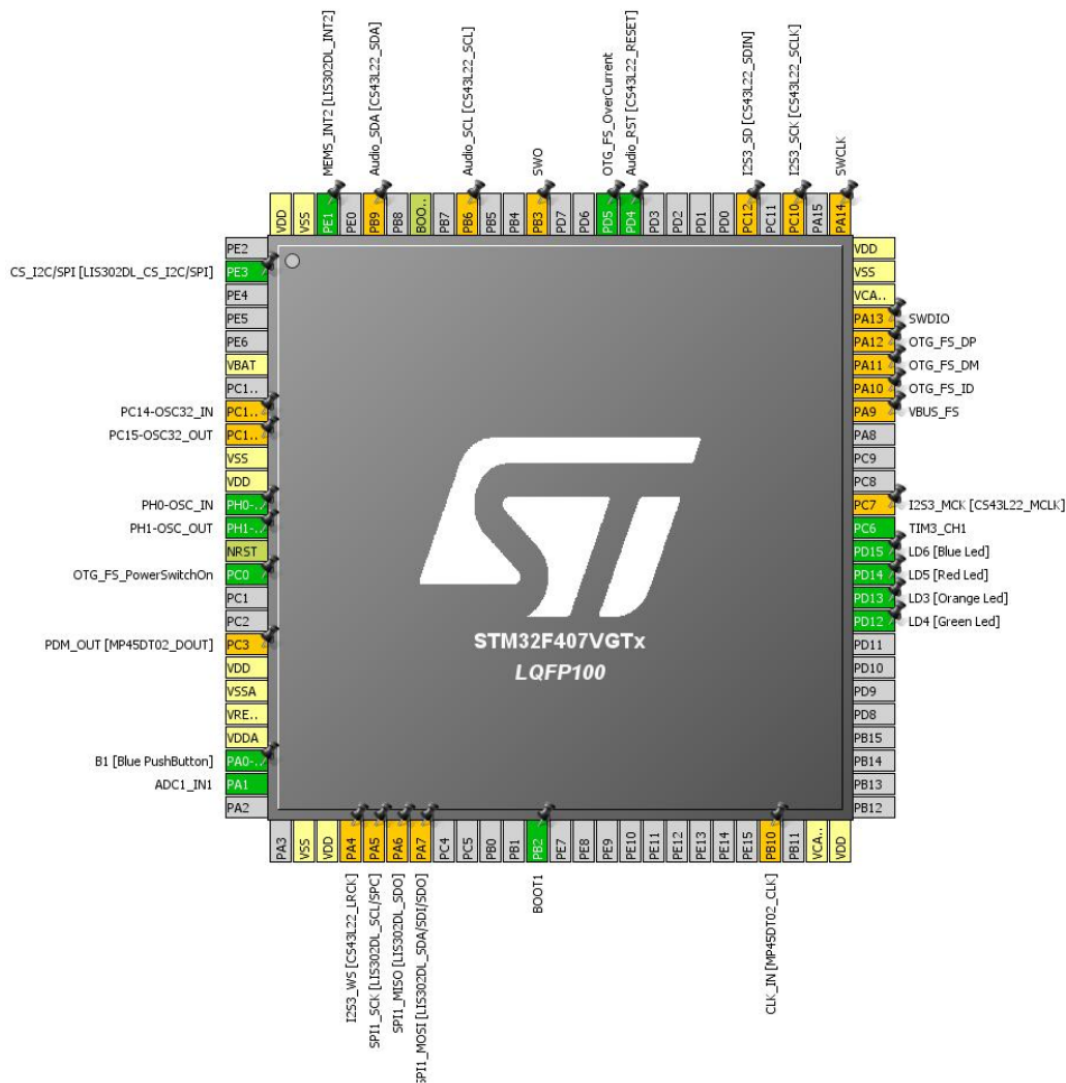


Figure F.1: HAL's Visualization of the STM32F4's Pins Configuration



# Appendix G

## Theory References

- 1. Tool, B. and Library, C. (2018). Rectifier Circuits — Diodes and Rectifiers — Electronics Textbook. [online] Allaboutcircuits.com. Available at: <https://www.allaboutcircuits.com/textbook/semiconductors/chapter-3/rectifier-circuits/> [Accessed 18 Mar. 2018].
- 2. Justsoftwaresolutions.co.uk. (2018). Locks, Mutexes, and Semaphores: Types of Synchronization Objects — Just Software Solutions - Custom Software Development. [online] Available at: <https://www.justsoftwaresolutions.co.uk/threading/locks-mutexes-semaphores.html> [Accessed 18 Mar. 2018].